

Lecture 7: Semaphores and monitors

601.418/618 Operating Systems

David Hovemeyer

February 12, 2025

Agenda

- ▶ Semaphores
- ▶ Monitors
- ▶ Condition variables

Acknowledgments: These slides are shamelessly adapted from [Prof. Ryan Huang's Fall 2022 slides](#), which in turn are based on [Prof. David Mazières's OS lecture notes](#).

Higher-Level Synchronization

Last time: We looked at using locks to provide mutual exclusion

Locks work, but they have limited semantics

- ▶ Just provide mutual exclusion

Instead, we want synchronization mechanisms that

- ▶ Block waiters, leave interrupts enabled in critical sections
- ▶ Provide semantics beyond mutual exclusion

Look at two common high-level mechanisms

- ▶ *Semaphores*: binary (mutex) and counting
- ▶ *Monitors*: mutexes and condition variables

Semaphores

An *abstract data type* to provide synchronization

- ▶ Described by Dijkstra in the “THE” system in 1968

Semaphores are “integers” that support two operations:

- ▶ Semaphore::P() decrements, blocks until semaphore is open, a.k.a wait()
 - ▶ after the Dutch word “Proberen” (to try)
 - ▶ Pintos `sema_down()`, pthreads `sem_wait()`
- ▶ Semaphore::V() increments, allows another thread to enter, a.k.a signal()
 - ▶ after the Dutch word “Verhogen” (increment)
 - ▶ Pintos `sema_up()`, pthreads `sem_post()`
- ▶ **That’s it!** No other operations – not even just reading its value

Semaphore safety property: the semaphore value is always greater than or equal to 0

Blocking in Semaphores

Associated with each semaphore is a queue of waiting threads

When $P()$ is called by a thread:

- ▶ If semaphore is *open*, thread continues
- ▶ If semaphore is *closed*, thread blocks on queue

Then $V()$ opens the semaphore:

- ▶ If a thread is waiting on the queue, the thread is unblocked
 - ▶ If multiple threads are waiting, one is chosen to wake up
- ▶ *If no threads are waiting on the queue, the signal is remembered for the next thread*
 - ▶ In other words, $V()$ has “history” (c.f., condition variables)
 - ▶ This “history” is a counter

Semaphore Types

Semaphores come in two types

Mutex semaphore (or *binary* semaphore)

- ▶ Represents single access to a resource
- ▶ Guarantees mutual exclusion to a critical section

Counting semaphore (or *general* semaphore)

- ▶ Represents a resource with many units available, or a resource that allows certain kinds of unsynchronized concurrent access (e.g., reading)
- ▶ Multiple threads can pass the semaphore
- ▶ Number of threads determined by the semaphore “count”
 - ▶ Mutex has initial count = 1, counting has count = N


Using Semaphores

Use is similar to locks (from last time), but semantics are different

```
struct Semaphore {  
    int value;  
    Queue q;  
} S;  
  
withdraw (account, amount) {  
    P(S);  
    balance = get_balance(account);  
    balance = balance - amount;  
    put_balance(account, balance);  
    v(S);  
    return balance;  
}
```

Threads
block
critical
section

It is undefined which
thread runs after a signal



Semaphore Questions

Are there any problems that can be solved with counting semaphores that cannot be solved with mutex semaphores?

- ▶ If a system only gives you mutex semaphore, can you use it to implement counting semaphores?

Does it matter which thread is unblocked by a signal operation?

Semaphore Implementation in Pintos

```
void sema_down(struct semaphore *sema)
{
    enum intr_level old_level;
    old_level = intr_disable();
    while (sema->value == 0) {
        list_push_back(&sema->waiters,
                       &thread_current()->elem);
        thread_block();
    }
    sema->value--;
    intr_set_level(old_level);
}
```

```
void sema_up(struct semaphore *sema)
{
    enum intr_level old_level;
    old_level = intr_disable();
    if (!list_empty (&sema->waiters))
        thread_unblock(list_entry(
            list_pop_front(&sema->waiters),
            struct thread, elem));
    sema->value++;
    intr_set_level(old_level);
}
```

To reference current thread: `thread_current()`

`thread_block()` puts the current thread to sleep

Assignment 1 note:

- leverage semaphore instead of directly using `thread_block()/thread_unblock()`

Implementation of thread_block()

pick another
thread to run



```
/* Puts the current thread to sleep. This function
must be called with interrupts turned off.*/
void thread_block ()
{
    ASSERT (!intr_context ());
    ASSERT (intr_get_level () == INTR_OFF);
    thread_current ()->status = THREAD_BLOCKED;
    schedule ();
}
```

thread_block() assumes the interrupts are disabled

This means we will have the thread sleep with interrupts disabled

Isn't this bad?

- ▶ Shouldn't we only disable interrupts when entering/leaving critical sections but keep interrupts enabled during critical section?

Interrupts Re-enabled Right After Context Switch

```
thread_yield() {  
    Disable interrupts;  
    add current thread to ready_list;  
    schedule(); // context switch  
    Enable interrupts;  
}
```

```
sema_down(){  
    Disable interrupts;  
    while(value == 0) {  
        add current thread to waiters;  
        thread_block();  
    }  
    value--;  
    Enable interrupts;  
}
```

[thread_yield]

```
Disable interrupts;  
add current thread to ready_list;  
schedule();
```

Thread 1

[thread_yield]

```
(Returns from schedule())  
Enable interrupts;
```

Thread 2

...

[sema_down]

```
Disable interrupts;  
while(value == 0) {  
    add current thread to waiters;  
    thread_block();  
}
```

Thread 2

[thread_yield]

```
(Returns from schedule())  
Enable interrupts;
```

Thread 1

Semaphore Summary

Semaphores can be used to solve any traditional synchronization problem

However, they have some drawbacks

- ▶ They are essentially shared global variables
 - ▶ Can potentially be accessed anywhere in program
- ▶ No connection between the semaphore and the data controlled by the semaphore
- ▶ Used both for critical sections (mutual exclusion) and coordination (scheduling)
 - ▶ Note that I had to use comments in the code to distinguish
- ▶ No control or guarantee of proper usage

Sometimes hard to use and prone to bugs

- ▶ Another approach: **Use programming language support**

Monitors

A programming language construct that controls access to shared data

- ▶ Synchronization code added by compiler, enforced at runtime
- ▶ **Why is this an advantage?**

A monitor is a module that encapsulates

- ▶ *Shared data structures*
- ▶ *Procedures* that operate on the shared data structures
- ▶ *Synchronization* between concurrent threads that invoke the procedures

A monitor protects its data from unstructured access

It guarantees that threads accessing its data through its procedures interact only in legitimate ways

Monitor Semantics

A monitor guarantees mutual exclusion

- ▶ Only one thread can execute any monitor procedure at any time
 - ▶ The thread is “in the monitor”
- ▶ If a second thread invokes a monitor procedure when a first thread is already executing one, it blocks
 - ▶ So the monitor has to have a wait queue. . .
- ▶ If a thread within a monitor blocks, another one can enter

What are the implications in terms of parallelism in a monitor?

A *monitor invariant* is a *safety property* associated with the monitor

- ▶ It's expressed over the monitored variables.
- ▶ It holds whenever a thread enters or exits the monitor.

Account Example

```
Monitor account {  
    double balance;  
  
    double withdraw(amount) {  
        balance = balance - amount;  
        return balance;  
    }  
}
```

Threads
block
waiting
to get
into
monitor

When first thread exits, another can
enter. Which one is undefined.

withdraw(amount)
balance = balance - amount;

withdraw(amount)

withdraw(amount)

return balance (and exit)

balance = balance - amount
return balance;

balance = balance - amount;
return balance;

Hey, that was easy!

Monitor invariant: $balance \geq 0$

Condition Variables

But what if a thread wants to wait for something inside the monitor?

- ▶ If we busy wait, it's bad
- ▶ Even worse, no one can get in the monitor to make changes now!

A *condition variable* is associated with a *condition* needed for a thread to make progress once it is in the monitor.

```
Monitor M {  
    ... monitored variables  
    Condition c;  
  
    void enterMonitor (...) {  
        if (extra property not true) wait(c);  waits outside of the monitor's mutex  
        do what you have to do  
        if (extra property true) signal(c);    brings in one thread waiting on condition  
    }  
}
```


Condition Variables

Condition variables support three operations:

- ▶ *Wait*: **release monitor lock**, wait for condition variable to be signaled
 - ▶ So condition variables have wait queues, too
- ▶ *Signal*: wake up one waiting thread
- ▶ *Broadcast*: wake up all waiting threads

Condition variables are not boolean objects

- ▶  `if (condition_variable) then ...` does not make sense
- ▶  `if (num_resources == 0) then wait(resources_available)` does
- ▶ An example later will make this more clear

Condition Vars $\neq$ Semaphores

Condition variables $\neq$ semaphores

- ▶ Although their operations have similar names, they have entirely different semantics (such is life, worse yet to come)
- ▶ **However, they each can be used to implement the other**

Access to the monitor is controlled by a lock

- ▶ `wait()` blocks the calling thread, and **gives up the lock**
 - ▶ To call `wait`, the thread has to be in the monitor (hence has lock)
 - ▶ `Semaphore::wait` just blocks the thread on the queue
- ▶ `signal()` causes a waiting thread to wake up
 - ▶ **If there is no waiting thread, the signal is lost**
 - ▶ `Semaphore::signal` increases the semaphore count, allowing future entry even if no thread is waiting
- ▶ Condition variables have no history

Signal Semantics

Two flavors of monitors that differ in the scheduling semantics of `signal()`

- ▶ Hoare monitors (original)
 - ▶ `signal()` **immediately switches from the caller to a waiting thread**
 - ▶ The condition that the waiter was anticipating is guaranteed to hold when waiter executes
 - ▶ Signaler must restore monitor invariants before signaling
- ▶ Mesa monitors (Mesa, Java)
 - ▶ `signal()` places a waiter on the ready queue, but **signaler continues inside monitor**
 - ▶ Condition is not necessarily true when waiter runs again
 - ▶ Returning from `wait()` is only a *hint* that something changed
 - ▶ Must recheck conditional case

Hoare vs. Mesa Monitors

Hoare monitor semantics:

```
if (!condition)
  wait(cond_var); // <-- condition definitely holds when wait() returns
```

Mesa/Java monitor semantics:

```
while (!condition)
  wait(cond_var); // <-- condition *might* hold when wait() returns
// <-- condition definitely holds when loop finishes
```

Tradeoffs:

- ▶ Mesa monitors easier to use, more efficient
 - ▶ Fewer context switches, easy to support broadcast
- ▶ Hoare monitors leave less to chance
 - ▶ Easier to reason about the program

Condition variables and locks

Condition variables are also used without monitors in conjunction with locks (e.g., pthreads)

```
void cond_init (cond_t *, ...);
```

```
void cond_wait (cond_t *c, mutex_t *m);
```

- ▶ Atomically unlock `m` and sleep until `c` signaled
- ▶ Then re-acquire `m` and resume executing

```
void cond_signal (cond_t *c);
```

```
void cond_broadcast (cond_t *c);
```

- ▶ Wake one/all threads waiting on `c`

Condition variables and locks

A monitor $\approx$ a module whose state includes condition variable(s) and a lock

- ▶ Difference is syntactic; with monitors, compiler adds the code

It is “just as if” each procedure in the module calls `acquire()` on entry and `release()` on exit

- ▶ But can be done anywhere in procedure, at finer granularity

With condition variables, the module methods may wait and signal on independent conditions

Condition variables and locks

Why must `cond_wait` both release `mutex_t` and block the caller?

► `void cond_wait(cond_t *c, mutex_t *m);`

Why not separate mutexes and condition variables?

```
while (count == BUFFER_SIZE) {  
    mutex_unlock(&mutex);  
    cond_wait(&not_full);  
    mutex_lock(&mutex);  
}
```

Condition variables and locks

Why must `cond_wait` both release `mutex_t` and block the caller?

► `void cond_wait(cond_t *c, mutex_t *m);`

Why not separate mutexes and condition variables?

Producer

```
while (count == BUFFER_SIZE) {  
    mutex_unlock(&mutex);  
  
    cond_wait(&not_full);  
    mutex_lock(&mutex);  
}
```

Consumer

```
mutex_lock(&mutex);  
... count--;  
cond_signal(&not_full);  
mutex_unlock(&mutex);
```


Using condition variables and locks

Alternation of two threads (ping-pong)

Each executes the following:

```
Lock lock;  
Condition cond;  
  
void ping_pong () {  
    acquire(lock);  
    while (1) {  
        printf("ping or pong\n");  
        signal(cond);  
        wait(cond, lock);  
    }  
    release(lock);  
}
```

Must acquire lock before you can wait
(similar to needing interrupts disabled
to call `thread_block` in Pintos)

Wait atomically releases lock
and blocks until `signal()`

After `signal()`, wait re-acquires
lock before returning

Monitors and Java

A lock and condition variable are in every Java object

- ▶ No explicit classes for locks or condition variables

Every object is/has a monitor

- ▶ At most one thread can be inside an object's monitor
- ▶ A thread enters an object's monitor by
 - ▶ Executing a method declared "synchronized"
 - ▶ Executing the body of a "synchronized" statement
- ▶ The compiler generates code to acquire the object's lock at the start of the method and release it just before returning
 - ▶ The lock itself is implicit, programmers do not worry about it

Monitors and Java

Every object can be treated as a condition variable

- ▶ Half of `Object`'s methods are for synchronization!

Take a look at the Java `Object` class:

- ▶ `Object.wait(*)` is `Condition::wait()`
- ▶ `Object.notify()` is `Condition::signal()`
- ▶ `Object.notifyAll()` is `Condition::broadcast()`

Summary

Semaphores

- ▶ `wait()/signal()` implement blocking mutual exclusion
- ▶ Also used as atomic counters (counting semaphores)
 - ▶ Often used to count availability of units of a resource
- ▶ Can be inconvenient to use

Monitors

- ▶ Synchronizes execution within procedures that manipulate encapsulated data shared among procedures
 - ▶ Only one thread can execute within a monitor at a time
- ▶ Relies upon high-level language support

Condition variables

- ▶ Used by threads as a synchronization point to wait for events
- ▶ Inside monitors, or outside with locks

Next time

Synchronization in practice