

Lecture 11: Virtual memory II

601.418/618 Operating Systems

David Hovemeyer

March 10, 2024

Agenda

- ▶ Hierarchical page tables
- ▶ TLB
- ▶ Handling page faults
- ▶ Memory-mapped files

Acknowledgments: These slides are shamelessly adapted from [Prof. Ryan Huang's Fall 2022 slides](#), which in turn are based on [Prof. David Mazières's OS lecture notes](#).

Lecture Overview

Today we'll cover more paging mechanisms:

Optimizations


- ▶ Managing page tables (space)
- ▶ Efficient translations (TLBs) (time)
- ▶ Demand paged virtual memory (space)

Recap address translation

Advanced Functionality

- ▶ Sharing memory
- ▶ Copy on Write
- ▶ Mapped files

Recap: Virtual Address Lookup in Page Table



Recap: Paging Example

Pages are 4K

- ▶ VPN is 20 bits (220 VPNs), offset is 12 bits


Virtual address is 0x7468

- ▶ Virtual page is 0x7, offset is 0x468

Page table entry 0x7 contains 0x2

- ▶ Physical page number is 0x2
- ▶ Seventh virtual page is at address 0x2000 (2nd physical page)

Physical address = $0x2000 + 0x468 = 0x2468$



Managing Page Tables

Size of the page table for a 32-bit address space w/ 4K pages

- ▶ $(2^{32}/2^{12}) \times 4\text{B} = 4\text{MB}$
- ▶ This is far far too much overhead for each process

How can we reduce this overhead?

- ▶ **Observation:** only need to map the portion of the address space actually being used (tiny fraction of entire address space)

How do we only map what is being used?

- ▶ Can dynamically extend page table...
- ▶ Does not work if address space is sparse (internal fragmentation)

Use another level of indirection: *two-level page tables*


Two-Level Page Tables

Two-level page tables

Virtual addresses (VAs) have **three** parts:

1. “Master” page number (index of entry in “root” page table)
2. Secondary page number (index of entry in “leaf” page table)
3. Offset (identifies byte in page)


Two-Level Page Tables



Two-Level Page Tables


Example

- ▶ 4KB-sized pages¹, 4 bytes/PTE
 - ▶ PTEs are generally the same size as an address: PPN is smaller than an address (omits offset bits), leaves room for metadata bits
- ▶ How many bits in offset? $\log_2(4K) = 12$ bits
- ▶ We want the master page table in one page: $4K/4\text{bytes} = 1K$ entries
- ▶ Hence, 1024 secondary page tables. How many bits?
- ▶ Master $\log_2(1K) = 10$, offset = 12, inner = $32 - 10 - 12 = 10$ bits



¹ $4K = 2^{12} = 4096$


x86 Page Translation




*32 bits aligned onto a 4-KByte boundary

Page Table Evolution



Linear (Flat)
Page Table





Virtual Address
Space



Physical Memory



Page Table Evolution



Wait a second...


We introduced two-level page tables to reduce the overhead of storing page tables

- ▶ Each flat page table costs $(2^{32}/2^{12}) \times 4\text{B} = 4\text{MB}$ to store

But even if we add another level, isn't the overhead the same?

- ▶ 1024 secondary page tables
- ▶ Each secondary page table has 2^{10} PTEs, thus has a size of 4KB
- ▶ Total size of these page tables is $1024 \times 4\text{KB} = 4\text{MB}$...
- ▶ In fact, we also have one master page table, which has a size of 4KB...

Page Table Evolution



Addressing Page Tables

Where do we store page tables (which address space)?

Physical memory

- ▶ Easy to address, no translation required
- ▶ But, allocated page tables consume memory for lifetime of VAS

Virtual memory (OS virtual address space)

- ▶ Cold (unused) page table pages can be paged out to disk
- ▶ But, addressing page tables requires translation
- ▶ How do we stop recursion?
- ▶ Do not page the outer page table (called *wiring*)

If we're going to page the page tables, might as well page the entire OS address space, too

- ▶ Need to wire special code and data (fault, interrupt handlers)

Efficient Translations

Our original page table already doubled the cost of memory access

- ▶ One lookup into the page table, another to fetch the data

Now two-level page tables triple the cost!

- ▶ Two lookups into the page tables, a **third** to fetch the data
- ▶ Worse, 64-bit architectures support 4-level page tables
- ▶ And this assumes the page table is in memory

How can we use paging but also reduce lookup cost?

- ▶ Cache translations in hardware
- ▶ Translation Lookaside Buffer (TLB)
- ▶ TLB managed by Memory Management Unit (MMU)

TLBs

Translation Lookaside Buffers

- ▶ Translate *virtual page #s* into *PTEs* (not physical addresses)
- ▶ Can be done in a single machine cycle


TLBs implemented in hardware

- ▶ Typically 4-way to fully associative cache (all entries looked up in parallel)
- ▶ Cache tags are virtual page numbers
- ▶ Cache values are PTEs (entries from page tables)
- ▶ With PTE + offset, can directly calculate physical address

TLBs exploit locality

- ▶ Processes only use a handful of pages at a time
 - ▶ 32-128 entries/pages (128-512K)
 - ▶ Only need those pages to be “mapped”
- ▶ Hit rates are therefore very important

TLBs



Managing TLBs

Address translations for most instructions are handled using the TLB

- ▶ > 99 % of translations, but there are misses (*TLB miss*)...

Who places translations into the TLB (loads the TLB)?

- ▶ *Hardware-managed TLB* (Memory Management Unit) [x86]
 - ▶ Knows where page tables are in main memory
 - ▶ OS maintains tables, HW accesses them directly
 - ▶ Tables have to be in HW-defined format (inflexible)
- ▶ *Software-managed TLB* (OS) [MIPS, Alpha, Sparc, PowerPC]
 - ▶ TLB faults to the OS, OS finds appropriate PTE, loads it in TLB
 - ▶ Must be fast (but still 20-200 cycles)
 - ▶ CPU ISA has instructions for manipulating TLB
 - ▶ Tables can be in any format convenient for OS (flexible)

Excellent overview of virtual memory and TLB management on modern(-ish) CPUs: B. Jacob and T. Mudge, [Virtual memory in contemporary microprocessors](#), IEEE Micro, July-August 1998

Managing TLBs (2)

OS ensures that TLB and page tables are consistent

- ▶ When it changes the protection bits of a PTE, it needs to invalidate the PTE if it is in the TLB

Reload TLB on a process context switch

- ▶ Invalidate all entries
- ▶ **Why? What is one way to fix it?**


When the TLB misses and a new PTE has to be loaded, a cached PTE must be evicted

- ▶ Choosing PTE to evict is called the *TLB replacement policy*
- ▶ Implemented in hardware, often simple (e.g., Not-Most-Recently-Used)

Paged Virtual Memory

Pages can be moved between memory and disk

- ▶ Use disk to simulate larger virtual than physical mem
- ▶ This process is called paging in/out



Paged Virtual Memory

Pages can be moved between memory and disk

Paging process over time

- ▶ Initially, pages are allocated from memory
- ▶ When memory fills up, allocating a page requires some other page to be evicted
- ▶ Evicted pages go to disk (**where? the swap file/backing store**)
- ▶ Done by the OS, and transparent to the application

Extreme design: demand paging

- ▶ Paging in a page from disk into memory only if an attempt is made to access it
- ▶ Main memory becomes a cache for disk

Page Faults


What happens when a process accesses a page is evicted?

1. When the OS evicts a page, it sets the PTE as invalid and stores the location of the page in the swap file in the PTE
2. When a process accesses the page, the invalid PTE causes a trap (*page fault*)
3. The trap will run the OS page fault handler
4. Handler uses the invalid PTE to locate page in swap file
5. Reads page into a physical frame, updates PTE to point to it
6. Restarts process

But where does it put it? Have to evict something else

- ▶ OS usually keeps a pool of free pages around so that allocations do not always cause evictions

Page Fault & Paging



Address Translation Redux

We started this topic with the high-level problem of translating virtual addresses into physical addresses

We've covered all of the pieces

- ▶ Virtual and physical addresses
- ▶ Virtual pages and physical page frames
- ▶ Page tables and page table entries (PTEs), protection
- ▶ TLBs
- ▶ Demand paging

Now let's put it together, bottom to top

Stay Cool



This is going to get a little bit involved.

The Common Case

Situation: Process is executing on the CPU, and it issues a read to an address

► **What kind of address is it? Virtual or physical?**

The read goes to the TLB in the MMU

1. TLB does a lookup using the *page number* of the address
2. Common case is that the page number matches, returning a *page table entry* (PTE) for the mapping for this address
3. TLB validates that the *PTE protection* allows reads (in this example)
4. PTE specifies which *physical frame* holds the page
5. MMU combines the physical frame and offset into a *physical address*
6. MMU then reads from that physical address, returns value to CPU

Note: **This is all done by the hardware**

TLB Misses, Protection Violations

At this point, two other things can happen

1. TLB does not have a PTE mapping this virtual address
2. PTE in TLB, but memory access violates PTE protection bits

We'll consider each in turn

Reloading the TLB

If the TLB does not have mapping, two possibilities:

1. MMU loads PTE from page table in memory
 - ▶ **Hardware managed TLB, OS not involved in this step**
 - ▶ OS has already set up the page tables so that the hardware can access it directly
2. Trap to the OS
 - ▶ **Software managed TLB, OS intervenes at this point**
 - ▶ OS does lookup in page table, loads PTE into TLB
 - ▶ OS returns from exception, TLB continues

A machine will only support one method or the other

At this point, there is a PTE for the address in the TLB

TLB Misses (2)

Note that:

Page table lookup (by HW or OS) can cause a recursive fault if page table is paged out

- ▶ Assuming page tables are in OS virtual address space
- ▶ Not a problem if tables are in physical memory
- ▶ Yes, this is a complicated situation

When TLB has PTE, it restarts translation

- ▶ Common case is that the PTE refers to a valid page in memory
 - ▶ These faults are handled quickly, just read PTE from the page table in memory and load into TLB
- ▶ Uncommon case is that TLB faults again on PTE because of PTE protection bits (e.g., page is invalid)
 - ▶ Becomes a page fault...

Page Faults


PTE can indicate a protection fault

- ▶ Read/write/execute – operation not permitted on page
- ▶ TLB traps to the OS (software takes over)
 - ▶ R/W/E – OS usually will send fault back up to process (e.g., segmentation fault), or might be playing games (e.g., copy on write, mapped files)

PTE could be marked as invalid

- ▶ Invalid = virtual page not allocated, or page not in physical memory
- ▶ Virtual page not allocated in address space
 - ▶ OS sends fault to process (e.g., segmentation fault)
- ▶ Page not in physical memory
 - ▶ OS allocates frame, reads from disk, maps PTE to physical frame

Address Translation: Putting It All Together



Advanced Functionality

Now we're going to look at some advanced functionality that the OS can provide applications using virtual memory tricks

- ▶ Shared memory
- ▶ Copy on Write
- ▶ Mapped files

Sharing

Private virtual address spaces protect applications from each other

- ▶ Usually exactly what we want

But this makes it difficult to share data (have to copy)

- ▶ Parents and children in a forking Web server or proxy will want to share an in-memory cache without copying

We can use *shared memory* to allow processes to share data using direct memory references

- ▶ Both processes see updates to the shared memory segment
 - ▶ Process B can immediately read an update by process A
- ▶ **How are we going to coordinate access to shared data?**

Sharing (2)

How can we implement sharing using page tables?

- ▶ Have PTEs in both tables map to the same physical frame²
- ▶ Each PTE can have different protection values
- ▶ Must update both PTEs when page becomes invalid


Can map shared memory at same or different virtual addresses in each process' address space

- ▶ Different: Flexible (no address space conflicts), but pointers inside the shared memory segment are invalid (**Why?**)
- ▶ Same: Less flexible, but shared pointers are valid (**Why?**)


What happens if a pointer inside the shared segment references an address outside the segment?

²Not directly possible on systems with “inverted” page tables, where physical pages are mapped back to virtual pages in a specific address space (e.g., PowerPC)

Isolation: No Sharing



Sharing Pages



Copy on Write


OSes spend a lot of time copying data

- ▶ System call arguments between user/kernel space
- ▶ Entire address spaces to implement `fork()`


Use Copy on Write (CoW) to defer large copies as long as possible, hoping to avoid them altogether

- ▶ Instead of copying pages, create *shared mappings* of parent pages in child virtual address space
- ▶ Shared pages are protected as read-only in parent and child
 - ▶ Reads happen as usual
 - ▶ Writes generate a protection fault, trap to OS, copy page, change page mapping in client page table, restart write instruction
- ▶ **How does this help `fork()`?**


Copy on Write: Before Fork



Copy on Write: Fork



Copy on Write: On a Write



Mapped Files

Mapped files enable processes to do file I/O using loads and stores

- ▶ Instead of “open, read into buffer, operate on buffer, ...”


Bind a file to a virtual memory region (`mmap()` in Unix)

- ▶ PTEs map virtual addresses to physical frames holding file data
- ▶ Virtual address base + N refers to offset N in file

Initially, all pages mapped to file are invalid

- ▶ OS reads a page from file when invalid page is accessed
- ▶ OS writes a page to file when evicted, or region unmapped
- ▶ If page is not dirty (has not been written to), no write needed
 - ▶ Another use of the dirty bit in PTE

Mapped Files



Mapped Files (2)

File is essentially backing store for that region of the virtual address space (instead of using the swap file)

- ▶ Virtual address space not backed by “real” files also called *Anonymous VM*

Advantages

- ▶ Uniform access for files and memory (just use pointers)
- ▶ Less copying (why?)

Drawbacks

- ▶ Process has less control over data movement
 - ▶ OS handles faults transparently
- ▶ Does not generalize to streamed I/O (pipes, sockets, etc.)

Summary

Paging mechanisms

Optimizations

- ▶ Managing page tables (space)
- ▶ Efficient translations (TLBs) (time)
- ▶ Demand paged virtual memory (space)

Recap address translation

Advanced Functionality

- ▶ Sharing memory
- ▶ Copy on Write
- ▶ Mapped files

Next time: Paging policies